

# Distributed PRNG Analysis Pipeline - Technical Addendum

Deeper details on feature contracts, sidecars, leakage guardrails, multi-model GPU behavior, and Step 6 pool scoring. Prepared by Team Beta - December 2025.

## What's inside

Section	What you'll learn
A1. Pipeline contract and artifacts	Exact handoffs, file shapes, and why sidecars exist.
A2. Feature system (48 per-seed + 14 more)	What "same features" means, and how ModelFactory uses the schema.
A3. Label leakage checks and why "RCentroid features" modes we already saw and how we guardrail them.	
A4. Multi-model training and GPU reality	My (OpenSource) CUDA exists and how it fits your modularity rules.
A5. Step 6 prediction pools and Jaccard	How intersection works and how to interpret agreement.
A6. Confidence scores: raw vs calibrated	Why you need "1.0" and how to use the three score types.
A7. Autonomy integration points to keep	Masters, compatibility, and lineage tracking.

## A1. Pipeline contract and artifacts

The system is designed around artifacts: outputs that can be validated, hashed, and consumed by later steps and automation. This prevents “magic state” and makes cluster execution reproducible.

### Key artifact contracts

Step	Primary script(s)	Produces	Consumed by
1	window_optimizer.py	optimal_window_config.json, bidirectional_scoring_worker.py	Step 2 / ScoringWatcherAgent evaluation
2/3	full_scoring_worker.py (+ job scripts)	survivors_with_scores.json (seed + features)	Step 5 (trainere)
4	meta_prediction_optimizer.py / adaptive_candidate_optimizer.py (historically)	best_model.* + best_model.meta.json	Step 5 (may ignore model-specific config w
5	meta_prediction_optimizer_anti_overfit.py	best_model.* + best_model.meta.json	Step 6 prediction_generator.py
6	prediction_generator.py	predictions_*.json (pool + scores)	WatcherAgent + downstream reporting

Rule (Team Beta): no file-extension guessing. Step 6 must load model\_type and checkpoint\_path only from best\_model.meta.json.

## A2. Feature system (48 per-seed + 14 global)

Your current training setup uses 48 per-seed features plus 14 global features (GlobalStateTracker). The per-seed features come from SurvivorScorer.extract\_ml\_features().

The trained model expects features in a fixed order. That order is declared in the sidecar (feature\_names) and validated by a feature\_schema\_hash. This prevents silent column shuffles.

### Are these the same 50 features from earlier?

Yes in spirit, with an important change: earlier files carried 50 keys inside the features dict, but two keys (score and confidence) are label-ish fields that can cause leakage. The current contract excludes them, leaving 48 true input features. Global features (14) are added at training and prediction time, giving 62 total inputs.

```
# Sidecar snippet (conceptual)
"feature_schema": {
    "per_seed_feature_count": 48,
    "global_feature_count": 14,
    "total_features": 62,
    "excluded_features": ["score", "confidence"],
    "ordering": "lexicographic_by_key",
    "feature_schema_hash": "..."
}
```

Guardrail: Step 5 must not re-run feature extraction when survivors\_with\_scores.json already contains precomputed features. Use the dicts as-is.

## A3. Label leakage checks and why “ $R^2 \approx 1$ ” can be a trap

A near-perfect  $R^2$  can be legitimate, but in this pipeline it is also a classic symptom of leakage: the label (or a proxy for it) sneaks into the features.

Guardrails you already applied: exclude score/confidence from features, log feature schema hash, and log label ranges (observed\_min/observed\_max).

Always validate that the trainer is learning from 48 per-seed feature inputs (plus 14 global), not from a hidden score field or from mutated arrays.

## A4. Multi-model training and GPU reality (OpenCL vs CUDA)

LightGBM often uses OpenCL on GPU, while PyTorch/XGBoost/CatBoost use CUDA. CUDA initialization can block OpenCL and cause LightGBM error -9999.

The robust fix is subprocess isolation: each trial runs in a fresh subprocess so GPU state is clean. The coordinator stays GPU-neutral and workers import GPU libraries only after parsing args.

```
# Coordinator pattern (conceptual)
# - main process: no torch/lightgbm imports
# - each trial: subprocess trains exactly one model
subprocess.run([
    'python3', 'train_single_trial.py',
    '--model-type', 'lightgbm',
    '--data-path', '/tmp/trial_data.npz',
    '--params', '{...}',
    '--save-model',
    '--model-output-dir', '/tmp/trial_models'
])
```

## A5. Step 6 prediction pools and Jaccard explained

Step 6 builds a pool of candidate next draws by scoring survivor candidates and aggregating their implied next outputs. Dual-sieve mode intersects forward and reverse survivors to focus on consistent candidates.

### Jaccard index

Given forward survivor set A and reverse survivor set B:

$$\text{Jaccard}(A, B) = |A \cap B| / |A \cup B|$$

Values near 1.0 mean strong agreement (intersection close to union). Values near 0 mean disagreement (small overlap). This is a great health metric for automation gating.

## A6. Confidence scores: raw vs calibrated vs normalized

Step 6 now emits three score tracks:

Field	What it is	Use it for
raw_scores	Direct model outputs per candidate	Automation thresholds; cross-run comparability
confidence_scores	Calibrated sigmoid(z-score) values	Readable certainty; avoids saturation
confidence_scores_normalized	Normalized-by-max display scores	Human UI only

Team Beta position: keep raw\_scores always; add normalization only as an extra field.

## A7. Autonomy integration points to keep Step 6 compatible

To keep the pipeline automation-ready, Step 6 must remain callable from CLI with stable argument names and must emit agent\_metadata consistent with the schema.

Item	Expectation
CLI args stable	No renames; add new flags with safe defaults (example: --parent-run-id optional)
Manifest updated	Step 6 manifest should declare new outputs (raw_scores, normalized)
Sidecar loading	best_model.meta.json is single source of truth for model_type + checkpoint_path + schema
Lineage tracking	Step 6 should accept --parent-run-id and also auto-read from sidecar as fallback

WatcherAgent decision policy is intentionally not specified here. This addendum focuses on strong contracts so a future policy can be added safely.